

# RAG Chatbot Assessment

## Assessment Overview

Develop a Retrieval-Augmented Generation (RAG) chatbot using FAISS as a vector database and LLM API (such as Gemini, OpenAI's GPT, or Hugging Face models), and LangChain's `create_retrieval_chain`. The goal is to create a chatbot that can provide accurate and contextually relevant responses to user queries.

## Key Focus Areas

- **RAG Workflow:** Efficient integration of retrieval-based search with response generation.
- **Vector Database:** FAISS for effective vector-based data storage and retrieval.
- **LLM API Integration:** Use a suitable LLM API, including options like Gemini, OpenAI, or Hugging Face models.
- **Document Loading:** Utilize document loaders for processing data sources (PDFs, Word documents).
- **API Design:** Build a user-interactive API endpoint for seamless querying.
- **LangChain Usage:** Leverage LangChain's `create_retrieval_chain` to handle document retrieval within the RAG process.

## Assessment Tasks

1. **Architecture Design**
  - **Task:** Design a high-level architecture for the RAG chatbot.
    - **Components:** Include an LLM (Gemini, OpenAI, or Hugging Face), FAISS vector database, LangChain's retrieval chain, and a chat API endpoint.
  - **Deliverable:** Submit a concise architecture overview.
2. **Document Loading and FAISS Vector Store Setup**
  - **Task:** Use LangChain's document loaders to load data from PDF or doc files and prepare them for retrieval with FAISS.
    - Load any publicly available documents (such as Wikipedia articles or research papers) to build data, and index this data in FAISS.
    - Configure retrieval to return the top 7 similar documents.
  - **Deliverable:** A FAISS-based vector store
3. **LLM API Integration with LangChain**
  - **Task:** Integrate an LLM API (Gemini, OpenAI, or Hugging Face models) using LangChain's `create_retrieval_chain`.
    - Ensure that FAISS retrieval provides relevant context to the LLM, which generates a response.

- Handle API authentication securely using environment variables (do not hardcode API keys).
  - **Deliverable:** Effective integration of the LLM API, ensuring accurate, context-based responses.
- 4. **Chat API Endpoint Implementation**
  - **Task:** Create a /chat endpoint that:
    - Accepts user questions, retrieves the top 7 relevant documents via FAISS, and generates a response using the LLM.
    - Returns a JSON response that includes the chatbot's answer, along with an optional list of source document filenames.
  - **Deliverable:** A functional and documented chat API endpoint.
- 5. **Optional Development (Encouraged but not required)**
  - **Memory Management:** Implement memory to retain conversation context across interactions.
  - **Source Tracking:** Include a list of source document filenames with each response.

### Additional Evaluation Criteria

- **Documentation:** Provide clear API documentation, including endpoint details, request/response formats, and sample calls.
- **Code Quality:** Ensure modular, readable code following Python best practices.
- **Efficiency:** Optimize for retrieval speed, response generation, and API performance.

### Submission Guidelines

- **Framework:** You may use any Python framework, such as Flask or FastAPI, for the implementation.
- **Repository/Zip File:** Push all code and documentation to a public GitHub repository (with at least 3–5 commits), and share the link with us. Alternatively, you can submit a zipped file of your project if preferred.
- **README:** Include a detailed README with setup instructions, environment variable usage, and guidance for running the application.
- **Timeline:** Complete the main tasks within 3 days. Share your progress by the deadline.

### Important Notes

- **Data Source Flexibility:** Use any publicly available documents for data retrieval.
- **API Key Security:** Use your own API key, but do not share it in the code or repository or with us. Use environment variables for secure handling.

FOR INTERNAL USE